

# COMP5112/MSBD5009 - MPI Manual

Please read the ***README.pdf*** first

## MPI environment setup

Please run `which mpicc` to check the MPI environment on the CS lab 2 workstations.

If you get outputs like:

```
mpicc: Command not found.
```

then you need to setup the MPI environment first:

- Add the OpenMPI installation path to your shell environment:

```
$ echo 'setenv PATH "${PATH}:/usr/local/software/openmpi/bin"' >> ~/.cshrc_user
```

- After that, re-login (logout and login) to active the new environment.
- You can use the command `which mpicc` to check your configuration.

## Compile and run MPI application

In this step, we take the `mpi_hello.c` as an example.

### Run MPI application on a single node

compile and run:

```
$ mpicc -std=c99 -o mpi_hello mpi_hello.c  
$ mpiexec -n <number of processes> ./mpi_hello
```

e.g. `mpiexec -n 2 ./mpi_hello`

If you need to run more processes than the physical cores of CPU, you need to add `--oversubscribe` to the `mpiexec` command:

e.g. `mpiexec --oversubscribe -n 8 ./mpi_hello`

An example output:

```
Greetings from process 0 of 8!  
Greetings from process 1 of 8!  
Greetings from process 2 of 8!  
Greetings from process 3 of 8!  
Greetings from process 4 of 8!  
Greetings from process 5 of 8!  
Greetings from process 6 of 8!  
Greetings from process 7 of 8!
```

## Run MPI application on multiple nodes

### Prepare hostfile

For running an MPI application on multiple nodes, we usually need to decide which nodes we will use to run the application.

We can put the hostnames or IPs into a hostfile.

For example, if we want to run our application on nodes: `cs12wk10`, `cs12wk11` and `cs12wk12`, the hostfile should be:

```
cs12wk10 slots=4  
cs12wk11 slots=4
```

```
csl2wk12 slots=4
```

Each node will have 4 available slots (12 slots in total).

The number of processes in the `mpiexec` command should be less than or equal to the total number of available slots (here is 12).

Prepare SSH

- check if you already have a RSA pri/pub key pair:

```
$ ls ~/.ssh
authorized_keys id_rsa id_rsa.pub known_hosts
```

if you cannot find `id_rsa` and `id_rsa.pub`, you need to generate your RSA keys first:

```
$ ssh-keygen -t rsa -b 4096
# repeatedly press <enter> until finish
```

- add your RSA key to `authorized_keys`:

```
$ touch ~/.ssh/authorized_keys
$ cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
```

- add target nodes to `known_hosts` (take `csl2wk10`, `csl2wk11` and `csl2wk12` as an example):

```
$ ssh csl2wk10
# type <yes> in your terminal
# you should login csl2wk10 successfully,
# please type <exit> in your terminal to back to your working workstation

$ ssh csl2wk11
# type <yes> in your terminal
# you should login csl2wk11 successfully,
# please type <exit> in your terminal to back to your working workstation

$ ssh csl2wk12
# type <yes> in your terminal
# you should login csl2wk12 successfully,
# please type <exit> in your terminal to back to your working workstation

# repeat above steps if you want to add more nodes to the `known_hosts`
```

compile and run your MPI application:

```
$ mpicc -std=c99 -o mpi_hello mpi_hello.c
$ mpiexec --hostfile hostfile -n <number of processes> ./mpi_hello
```

e.g. `mpiexec --hostfile hostfile -n 12 ./mpi_hello`

An example output:

```
Greetings from process 0 of 12!
Greetings from process 1 of 12!
Greetings from process 2 of 12!
Greetings from process 3 of 12!
Greetings from process 4 of 12!
Greetings from process 5 of 12!
Greetings from process 6 of 12!
Greetings from process 7 of 12!
Greetings from process 8 of 12!
Greetings from process 9 of 12!
```

